

# Project Mercedes-Benz Greener Manufacturing

August 19, 2022

*Project Mercedes-Benz Greener Manufacturing*

```
[1]: #Importing the necessary library
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[2]: #Getting the data and storing it in a DataFrame
merc_data_df=pd.read_csv('train.csv')
merc_data_df
```

```
[2]:
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	\
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
4204	8405	107.39	ak	s	as	c	d	aa	d	q	...	1	0	0	0	
4205	8406	108.77	j	o	t	d	d	aa	h	h	...	0	1	0	0	
4206	8412	109.22	ak	v	r	a	d	aa	g	e	...	0	0	1	0	
4207	8415	87.48	al	r	e	f	d	aa	l	u	...	0	0	0	0	
4208	8417	110.85	z	r	ae	c	d	aa	g	w	...	1	0	0	0	

	X379	X380	X382	X383	X384	X385
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	1	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0
...	...	...	...	...	...	...
4204	0	0	0	0	0	0
4205	0	0	0	0	0	0
4206	0	0	0	0	0	0
4207	0	0	0	0	0	0
4208	0	0	0	0	0	0

[4209 rows x 378 columns]

```
[3]: merc_data_df.shape
```

```
[3]: (4209, 378)
```

```
[4]: merc_data_df['y'].value_counts()
```

```
[4]: 90.76      7
      89.06      7
      89.38      7
      91.88      7
      93.62      6
      ..
      93.26      1
      93.24      1
      105.94      1
      94.17      1
      79.00      1
      Name: y, Length: 2545, dtype: int64
```

```
[5]: #checking for the Variance of each and every column
merc_var=pd.DataFrame(merc_data_df.var())
merc_var.columns=['Variance']
print(merc_var)
print(merc_var.loc[merc_var['Variance']==0]) #This gives me the columns with
↳ Variance value equal to 0
#method2_to_get_list_of_columns_with_variance_zero
#var_zero_list=merc_var[merc_var==0].dropna() #Gives Nan value to all the
↳ columns with a variance other the zero but in our case since we need only
↳ the ones with 0 variance, so we have removed all the Nan columns using dropna
#print(var_zero_list.loc)
```

	Variance
ID	5.941936e+06
y	1.607667e+02
X10	1.313092e-02
X11	0.000000e+00
X12	6.945713e-02
...	...
X380	8.014579e-03
X382	7.546747e-03
X383	1.660732e-03
X384	4.750593e-04
X385	1.423823e-03

```
[370 rows x 1 columns]
```

	Variance
X11	0.0

```

X93      0.0
X107     0.0
X233     0.0
X235     0.0
X268     0.0
X289     0.0
X290     0.0
X293     0.0
X297     0.0
X330     0.0
X347     0.0

```

```
[6]: merc_data_df['X11']
```

```

[6]: 0      0
     1      0
     2      0
     3      0
     4      0
     ..
    4204    0
    4205    0
    4206    0
    4207    0
    4208    0
     Name: X11, Length: 4209, dtype: int64

```

*#Here i have checked a Variable with Zero Variance and as you can see it is a single deterministic value. This might be a issue when we are running certain algorithms*

```

[7]: #Dropping the Columns with Zero Variance
merc_data_df.
      ↪drop(['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290', 'X293', 'X297', 'X330', 'X347'],ax
merc_data_df.head()

```

```

[7]:   ID      y  X0 X1  X2 X3 X4 X5 X6 X8  ...  X375  X376  X377  X378  X379  \
0   0  130.81   k  v  at  a  d  u  j  o  ...    0    0    1    0    0
1   6   88.53   k  t  av  e  d  y  l  o  ...    1    0    0    0    0
2   7   76.26  az  w   n  c  d  x  j  x  ...    0    0    0    0    0
3   9   80.62  az  t   n  f  d  x  l  e  ...    0    0    0    0    0
4  13   78.02  az  v   n  f  d  h  d  n  ...    0    0    0    0    0

      X380  X382  X383  X384  X385
0      0      0      0      0      0
1      0      0      0      0      0
2      0      1      0      0      0
3      0      0      0      0      0
4      0      0      0      0      0

```

[5 rows x 366 columns]

*Removing the columns with Zero Variance cause they have a effect on PCA algorithm. Since this dataset is having large no of features it is better to perform PCA and bring down no of variables. Since the column with Zero Variance is a blocker for PCA we are removing those columns from our Dataset*

```
[8]: merc_data_df.describe()
```

```
[8]:
```

	ID	y	X10	X12	X13	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	
mean	4205.960798	100.669318	0.013305	0.075077	0.057971	
std	2437.608688	12.679381	0.114590	0.263547	0.233716	
min	0.000000	72.110000	0.000000	0.000000	0.000000	
25%	2095.000000	90.820000	0.000000	0.000000	0.000000	
50%	4220.000000	99.150000	0.000000	0.000000	0.000000	
75%	6314.000000	109.010000	0.000000	0.000000	0.000000	
max	8417.000000	265.320000	1.000000	1.000000	1.000000	

  

	X14	X15	X16	X17	X18	...	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	...	
mean	0.428130	0.000475	0.002613	0.007603	0.007840	...	
std	0.494867	0.021796	0.051061	0.086872	0.088208	...	
min	0.000000	0.000000	0.000000	0.000000	0.000000	...	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	...	
75%	1.000000	0.000000	0.000000	0.000000	0.000000	...	
max	1.000000	1.000000	1.000000	1.000000	1.000000	...	

  

	X375	X376	X377	X378	X379	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000	
mean	0.318841	0.057258	0.314802	0.020670	0.009503	
std	0.466082	0.232363	0.464492	0.142294	0.097033	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	1.000000	0.000000	1.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	1.000000	

  

	X380	X382	X383	X384	X385
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
mean	0.008078	0.007603	0.001663	0.000475	0.001426
std	0.089524	0.086872	0.040752	0.021796	0.037734
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000

75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

[8 rows x 358 columns]

*From this also we can infer if a column has zero variance or not just by looking into the row of Standard Deviation*

```
[9]: #Importing the Test set Data
merc_test_df=pd.read_csv('test.csv')
merc_test_df.head()
```

```
[9]:   ID  X0 X1  X2 X3 X4 X5 X6 X8  X10  ...  X375  X376  X377  X378  X379  X380  \
0    1  az  v   n  f  d  t  a  w    0  ...    0    0    0    1    0    0
1    2   t  b  ai  a  d  b  g  y    0  ...    0    0    1    0    0    0
2    3  az  v  as  f  d  a  j  j    0  ...    0    0    0    1    0    0
3    4  az  l   n  f  d  z  l  n    0  ...    0    0    0    1    0    0
4    5   w  s  as  c  d  y  i  m    0  ...    1    0    0    0    0    0

      X382  X383  X384  X385
0         0         0         0         0
1         0         0         0         0
2         0         0         0         0
3         0         0         0         0
4         0         0         0         0
```

[5 rows x 377 columns]

```
[10]: #checking for the Variance of each and every column
merc_test_var=pd.DataFrame(merc_test_df.var())
merc_test_var.columns=['Variance']
print(merc_test_var)
print(merc_test_var.loc[merc_test_var['Variance']==0]) #This gives me the
↳ columns with Variance value equal to 0
```

	Variance
ID	5.871311e+06
X10	1.865006e-02
X11	2.375861e-04
X12	6.885074e-02
X13	5.734498e-02
...	...
X380	8.014579e-03
X382	8.715481e-03
X383	4.750593e-04
X384	7.124196e-04
X385	1.660732e-03

[369 rows x 1 columns]

	Variance
X257	0.0
X258	0.0
X295	0.0
X296	0.0
X369	0.0

```
[11]: #Dropping the Columns with Zero Variance
merc_test_df.drop(['X257', 'X258', 'X295', 'X296', 'X369'],axis=1,inplace=True)
merc_test_df.head()
```

```
[11]:   ID  X0 X1  X2 X3 X4 X5 X6 X8  X10 ...  X375  X376  X377  X378  X379  X380  \
0   1  az  v   n  f  d  t  a  w    0 ...    0    0    0    1    0    0
1   2   t  b  ai  a  d  b  g  y    0 ...    0    0    1    0    0    0
2   3  az  v  as  f  d  a  j  j    0 ...    0    0    0    1    0    0
3   4  az  l   n  f  d  z  l  n    0 ...    0    0    0    1    0    0
4   5   w  s  as  c  d  y  i  m    0 ...    1    0    0    0    0    0

      X382  X383  X384  X385
0         0         0         0         0
1         0         0         0         0
2         0         0         0         0
3         0         0         0         0
4         0         0         0         0
```

[5 rows x 372 columns]

*Since we are going to do PCA to Test Data As well i might throw an error if we have columns with zero variance, So i am removing it from the Test Dataset*

```
[12]: #Checking for null values in both the Train Dataset
merc_train_is_null=pd.DataFrame(merc_data_df.isnull().sum())
merc_train_is_null.columns=['Total_Null_values']
print(merc_train_is_null)
print(merc_train_is_null.loc[merc_train_is_null['Total_Null_values']!=0])
```

	Total_Null_values
ID	0
y	0
X0	0
X1	0
X2	0
...	...
X380	0
X382	0
X383	0
X384	0

X385                      0

```
[366 rows x 1 columns]
Empty DataFrame
Columns: [Total_Null_values]
Index: []
```

*From the last print Statement we can infer that no column has null value, here i basically checked if any column had a null value. Usually the first line alon is sufficient to infer but since here it is not displaying all the the columns, so i had to check it via code.*

```
[13]: #Checking for null values in both the Test Dataset
merc_test_is_null=pd.DataFrame(merc_test_df.isnull().sum())
merc_test_is_null.columns=['Total_Null_values']
print(merc_test_is_null)
print(merc_test_is_null.loc[merc_test_is_null['Total_Null_values']!=0])
```

	Total_Null_values
ID	0
X0	0
X1	0
X2	0
X3	0
...	...
X380	0
X382	0
X383	0
X384	0
X385	0

```
[372 rows x 1 columns]
Empty DataFrame
Columns: [Total_Null_values]
Index: []
```

*From the last print Statement we can infer that no column has null value, here i basically checked if any column had a null value.*

```
[14]: #Checking uniques Values of Train Dataset
print(merc_data_df['X0'].value_counts())
print(merc_data_df.nunique())
```

z	360
ak	349
y	324
ay	313
t	306
x	300
o	269

f	227
n	195
w	182
j	181
az	175
aj	151
s	106
ap	103
h	75
d	73
al	67
v	36
af	35
ai	34
m	34
e	32
ba	27
at	25
a	21
ax	19
i	18
aq	18
am	18
u	17
aw	16
l	16
ad	14
b	11
au	11
k	11
r	10
as	10
bc	6
ao	4
c	3
aa	2
q	2
ab	1
ac	1
g	1
Name: X0, dtype: int64	
ID	4209
y	2545
X0	47
X1	27
X2	44
...	
X380	2



```
X382      2
X383      2
X384      2
X385      2
Length: 366, dtype: int64
```

*From this i can infer how many unique values are available in each column of Dataset*

```
[15]: #Checking uniques Values of Test Dataset
print(merc_test_df['X0'].value_counts())
print(merc_test_df.nunique())
```

```
ak      432
y       348
z       335
x       302
ay      299
t       293
o       246
f       213
w       198
j       171
n       167
aj      162
az      161
s       116
ap      108
al       88
h        64
d        61
e        48
v        40
ai       38
m        34
af       34
am       28
i        25
at       21
u        20
ba       19
a        18
b        13
k        12
ad       12
aq       11
aw       11
r        10
ax         8
```

```

bc      6
l       6
as      6
c       6
au      5
ao      5
g       3
an      1
av      1
ae      1
bb      1
ag      1
p       1
Name: X0, dtype: int64
ID      4209
X0      49
X1      27
X2      45
X3       7
...
X380    2
X382    2
X383    2
X384    2
X385    2
Length: 372, dtype: int64

```

*#Setting Up Data for Model Building*

```

[16]: #splitting the train dataset into input and output for the model
x_train=merc_data_df.drop('y',axis=1)
print(x_train.head())
y_train=merc_data_df['y']
print(y_train.head())

```

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	\
0	0	k	v	at	a	d	u	j	o	0	...	0	0	1	0	0	0	
1	6	k	t	av	e	d	y	l	o	0	...	1	0	0	0	0	0	
2	7	az	w	n	c	d	x	j	x	0	...	0	0	0	0	0	0	
3	9	az	t	n	f	d	x	l	e	0	...	0	0	0	0	0	0	
4	13	az	v	n	f	d	h	d	n	0	...	0	0	0	0	0	0	

  

	X382	X383	X384	X385
0	0	0	0	0
1	0	0	0	0
2	1	0	0	0
3	0	0	0	0
4	0	0	0	0

```
[5 rows x 365 columns]
0    130.81
1     88.53
2     76.26
3     80.62
4     78.02
Name: y, dtype: float64
```

```
[17]: #Assigning Test Data to a variable
x_test=merc_test_df
print(x_test.head())
```

```

   ID  X0 X1  X2 X3 X4 X5 X6 X8  X10  ...  X375  X376  X377  X378  X379  X380  \
0   1  az  v   n  f  d  t  a  w    0  ...    0     0     0     1     0     0
1   2   t  b  ai  a  d  b  g  y    0  ...    0     0     1     0     0     0
2   3  az  v  as  f  d  a  j  j    0  ...    0     0     0     1     0     0
3   4  az  l   n  f  d  z  l  n    0  ...    0     0     0     1     0     0
4   5   w  s  as  c  d  y  i  m    0  ...    1     0     0     0     0     0

   X382  X383  X384  X385
0      0      0      0      0
1      0      0      0      0
2      0      0      0      0
3      0      0      0      0
4      0      0      0      0
```

```
[5 rows x 372 columns]
```

```
[18]: x_test.head()
```

```

[18]:   ID  X0 X1  X2 X3 X4 X5 X6 X8  X10  ...  X375  X376  X377  X378  X379  X380  \
0   1  az  v   n  f  d  t  a  w    0  ...    0     0     0     1     0     0
1   2   t  b  ai  a  d  b  g  y    0  ...    0     0     1     0     0     0
2   3  az  v  as  f  d  a  j  j    0  ...    0     0     0     1     0     0
3   4  az  l   n  f  d  z  l  n    0  ...    0     0     0     1     0     0
4   5   w  s  as  c  d  y  i  m    0  ...    1     0     0     0     0     0

   X382  X383  X384  X385
0      0      0      0      0
1      0      0      0      0
2      0      0      0      0
3      0      0      0      0
4      0      0      0      0
```

```
[5 rows x 372 columns]
```

*Since both the Test and Train Data have certain columns with Object Datatype, it is better to convert*

them into numerical datatype for model training

```
[19]: #Label_Encoding on Train_Data
from sklearn.preprocessing import LabelEncoder
lab_enc=LabelEncoder()
train_columns=x_train.columns    #To get the column Names into a list so that i
    ↪ can use it to loop and fit on every column
print(train_columns)
for i in train_columns:
    lab_enc.fit(x_train[i])      #Train on the data
    x_train[i]=lab_enc.transform(x_train[i]) #transforming the data
print(x_train.head())
```

```
Index(['ID', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8', 'X10',
      ...
      'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
      'X385'],
      dtype='object', length=365)
   ID  X0  X1  X2  X3  X4  X5  X6  X8  X10  ...  X375  X376  X377  X378  X379  \
0   0   32  23  17   0   3  24   9  14   0  ...    0     0     1     0     0
1   1   32  21  19   4   3  28  11  14   0  ...    1     0     0     0     0
2   2   20  24  34   2   3  27   9  23   0  ...    0     0     0     0     0
3   3   20  21  34   5   3  27  11   4   0  ...    0     0     0     0     0
4   4   20  23  34   5   3  12   3  13   0  ...    0     0     0     0     0

   X380  X382  X383  X384  X385
0     0     0     0     0     0
1     0     0     0     0     0
2     0     1     0     0     0
3     0     0     0     0     0
4     0     0     0     0     0
```

[5 rows x 365 columns]

```
[20]: #Label_Encoding on Test_data
from sklearn.preprocessing import LabelEncoder
lab_enc_test=LabelEncoder()
test_columns=x_test.columns    #To get the column Names into a list so that i
    ↪ can use it to loop and fit on every column
print(test_columns)
for j in test_columns:
    lab_enc_test.fit(x_test[j]) #Train on the data
    x_test[j]=lab_enc_test.transform(x_test[j]) #transforming the data
print(x_test.head())
```

```
Index(['ID', 'X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8', 'X10',
      ...
      'X375', 'X376', 'X377', 'X378', 'X379', 'X380', 'X382', 'X383', 'X384',
```

```

        'X385'],
        dtype='object', length=372)
    ID  X0  X1  X2  X3  X4  X5  X6  X8  X10  ...  X375  X376  X377  X378  X379  \
0    0   21  23  34   5   3  26   0  22   0  ...    0     0     0     1     0
1    1   42   3   8   0   3   9   6  24   0  ...    0     0     1     0     0
2    2   21  23  17   5   3   0   9   9   0  ...    0     0     0     1     0
3    3   21  13  34   5   3  31  11  13   0  ...    0     0     0     1     0
4    4   45  20  17   2   3  30   8  12   0  ...    1     0     0     0     0

        X380  X382  X383  X384  X385
0          0     0     0     0     0
1          0     0     0     0     0
2          0     0     0     0     0
3          0     0     0     0     0
4          0     0     0     0     0

```

[5 rows x 372 columns]

*#The major takeaways from the Label Encoding is that actually we must fit on the train set and then transform both train and test set. But here due improper dataset if i try to train on train dataset and transform both datasets, it is throwing an error because if you observe carefully, we can infer that Train Data has less no of features when compared to Test Set, so, if we try to transform the test set on the basis of fitted train set. Then it will throw errors cause some features are missing in test set. SO, we can also infer that it is better to do dimensionality reduction and bring down no of features similar in both Test and Train Datasets.*

*#Dimensionality Reduction*

```

[21]: #Dimensionality Reduction using Principal Component Analysis for Train Set
from sklearn.decomposition import PCA
pca_merc_train=PCA(n_components=10) #here n_components is hyperparameter
pca_merc_train.fit(x_train) #training on the x_train Data
print(pca_merc_train.explained_variance_ratio_) #Inorder to check the Total
→Variation our algo is accounting to
x_train_trans=pca_merc_train.transform(x_train) #Transforming our training
→data, that is bringing down the no of features

```

```

[9.99659608e-01 1.38083753e-04 7.69770342e-05 4.40107808e-05
 3.31849481e-05 2.66009200e-05 5.73061934e-06 2.67845282e-06
 1.56203752e-06 1.05802559e-06]

```

```

[22]: #Dimensionality Reduction using Principal Component Analysis for Test Set
from sklearn.decomposition import PCA
pca_merc_test=PCA(n_components=10) #here n_components is hyperparameter
pca_merc_test.fit(x_test) #training on the x_test Data
print(pca_merc_test.explained_variance_ratio_) #Inorder to check the Total
→Variation our algo is accounting to

```

```
x_test_trans=pca_merc_test.transform(x_test) #Transforming our testing data,
↳that is bringing down the no of features
print(x_test_trans)
```

```
[9.99637654e-01 1.67216554e-04 6.79028030e-05 4.33597184e-05
 3.31749814e-05 2.93170202e-05 5.51112705e-06 2.77341533e-06
 1.56251928e-06 1.04907502e-06]
[[ 2.10391944e+03  1.48901578e+01  1.43106871e+01 ... -1.71520611e+00
  2.68015002e+00 -1.39321798e+00]
 [ 2.10300998e+03 -1.48068922e+01 -8.10199579e+00 ...  4.15276621e+00
  1.92372374e+00  4.46863094e-01]
 [ 2.10203711e+03  1.23521650e+01 -2.18928463e+00 ... -8.20388024e-01
  7.81621434e-01 -1.89336891e-01]
...
 [-2.10191831e+03 -1.37536955e+01  3.21752326e+00 ... -2.88530584e+00
 -1.06187275e+00 -2.33170640e+00]
 [-2.10292321e+03  2.46245283e+01 -5.02830129e+00 ...  3.48120149e+00
 -2.02639839e+00  3.81031514e-02]
 [-2.10390525e+03 -1.56684906e+01 -7.94246831e+00 ...  7.06982115e-02
  1.06468235e+00  2.08030214e+00]]
```

*#Main takeaways from Principal Component Analysis is that, actually fitting should only be performed on train data and using this we need to transform our train and test data, but here since the train and test data are from different datasets and also they have unequal number of features, it will throw an size mismatch error when we try to transform of test data. So, that is the reason why i fitted on the test data as well. Following this, since n\_components is a hyperparameter it is difficult to guess the exact number to use. Here we transform our data which had around 350-380 features to a dataset which has only 10-30 features.*

```
[23]: #Building XGBoost Model
import xgboost
xgb_regressor=xgboost.XGBRegressor() #Since the variable to be predicted is
↳Continous we are going to using regression algo here
```

*#Here since we dont know the optimal parameter values to be used, it is better to use Gridsearch and find the right parameters*

```
[24]: #checking the GridSearchCV Algorithm to find the best parameter values
from sklearn.model_selection import GridSearchCV
params={'n_estimators':[100, 200, 400, 800], 'max_depth':[1,2,3,6,10]}
grid_search_cv=GridSearchCV(xgb_regressor,params,cv=3,n_jobs=-1)
grid_search_cv.fit(x_train_trans,y_train)
print(grid_search_cv.best_params_)
```

```
{'max_depth': 3, 'n_estimators': 100}
```

```
[25]: #Using Gridsearch Algo to find the best parameter value of learning rate and
↳min child weight
```

```

params_2={'learning_rate' : [0.1, 0.2, 0.3, 0.5], 'min_child_weight' : [1, 2, 3, 4, 5]}
grid_search_cv_2=GridSearchCV(xgb_regressor,params_2,cv=3,n_jobs=-1)
grid_search_cv_2.fit(x_train_trans,y_train)
print(grid_search_cv_2.best_params_)

```

```

{'learning_rate': 0.1, 'min_child_weight': 4}

```

```

[26]: #Using Gridsearch Algo to find the best parameter value of subsample
params_3={'subsample' : [0.5, 0.6, 0.7, 0.8, 1.0]}
grid_search_cv_3=GridSearchCV(xgb_regressor,params_3,cv=3,n_jobs=-1)
grid_search_cv_3.fit(x_train_trans,y_train)
print(grid_search_cv_3.best_params_)

```

```

{'subsample': 1.0}

```

*#From the above part of the code we can infer that we are tuning to find the best parameters to run the xgbregressor algorithm. Here first i tried tuning all the 5 parameters together but it requires a lot of time since here we were trying to replicate almost 100 regressor DT with different parameters which leads to a lot of combination of trees. So to reduce this i have splitted them into different sets and tunned to find the best parameters. 1)started with tuning n\_estimators which gives the number of boosting rounds required or number of trees to built. 2)max\_depth: maximum tree depth required for baselearners, it should be in a optimal window, if it is too high the tree becomes more complex and tends to overfit 3)Learning rate: This also determines how our model coverges, if too high it gets difficult to converge, if low it might take a lot of boosting rounds to converge. 4)min\_child\_weight: gives the number of child nodes required at the present node. 5)subsample: subsampling ratio of the training instances. It will occur once in every boosting iteration. Subsample ratio = 0.5 means that the algorithm would randomly sample half of the training data prior to growing trees.*

```

[27]: #Building the tree with optimal parameters for improved prediction
xgb_regressor=xgboost.
    ↳XGBRFRegressor(n_estimators=100,max_depth=3,learning_rate=0.
    ↳1,min_child_weight=5,subsample=1.0)
xgb_regressor.fit(x_train_trans,y_train)    #Training the model

```

```

[27]: XGBRFRegressor(base_score=0.5, booster=None, colsample_bylevel=1,
                    colsample_bytrees=1, gamma=0, gpu_id=-1, importance_type='gain',
                    interaction_constraints=None, learning_rate=0.1,
                    max_delta_step=0, max_depth=3, min_child_weight=5, missing=nan,
                    monotone_constraints=None, n_estimators=100, n_jobs=0,
                    num_parallel_tree=100, objective='reg:squarederror',
                    random_state=0, reg_alpha=0, scale_pos_weight=1, subsample=1.0,
                    tree_method=None, validate_parameters=False, verbosity=None)

```

```

[28]: #Preditcing the Test_df Values
y_test=xgb_regressor.predict(x_test_trans)
print(y_test)

```

```
[ 9.334505  9.703091 10.05097 ... 10.969909 11.347246 10.152216]
```

*#Finally, to conclude this we have first checked and removed any columns with zero variance just to overcome the situation where PCA algorithm might throw an error. Then we did dimensionality reduction using PCA to bring down the number of features. Then we built a model to predict the test\_df values using the test data. Here since the test data doesnot have the output variable it is not possible to get the mean squared error for this dataset*